k-means Is All You Need

Leonardo Biason (2045751)   Alessandro Romania (2046144)   Davide De Blasio (2082600)

Abstract—The k-means algorithm is a well known clustering algorithm, which is often used in unsupervised learning settings. However, the algorithm requires to perform multiple times the same operation on the data, and it can greatly benefit from a parallel implementation, so that to maximize the throughput and reduce computation times. With this project, we propose some possible implementations, based on some libraries that are considered to be the de-facto standard when it comes to writing multithreaded or parallel code, and we will discuss also the results of such implementations

Sapienza, ACSAI, Multicore Programming

I. INTRODUCTION

When talking about clustering and unsupervised learning, it's quite common to hear about the k-means algorithm, and for good reasons: it allows to efficiently cluster a dataset of d dimensions, and it employs the notion of convergence in order to do so. This, computationally speaking, means to repeat some operations over and over again until some stopping conditions are met.

The algorithm is not perfect though, and presents some issues:

1) the algorithm is fast in clustering, but we cannot be certain that it clusters *well*;
2) the algorithm doesn't work with non-linear clusters;
3) the initialization can make a great impact in the final result.

Many people prefer to use other clustering methods, such as the fitting of Gaussian Mixture Models. Albeit not being perfect, k-means still works well in simple, linear clusters. For the sake of this project, we are going to consider a vanilla k-means algorithm with Lloyd's initialization (the first k centroids will be selected randomly).

A. Algorithm structure

The k-means algorithm can be described with the following pseudocode, where X is the set of data points, $C = \{\mu_1, \mu_2, ..., \mu_k\}$ is the set of centroids and Y is the set of assignments:

Algorithm 1: k-means (Lloyd's initialization)

   // Initialize the centroids
1 **for** k *in* $[1, |C|]$ **do**
2     $\mu_k \leftarrow$ a random location in the input space
3 **end**
4 **while** *convergence hasn't been reached* **do**
     // Assign each point to a cluster
5   **for** i *in* $[1, |X|]$ **do**
6      $y_i \leftarrow \mathrm{argmin}_k \left(\|\mu_k - x_i\| \right)$
7   **end**
     // Compute the new position of each centroid
8   **for** k *in* $[1, |C|]$ **do**
9      $\mu_k \leftarrow \mathrm{MEAN}(\{ x_n : z_n = k \})$
10   **end**
11 **end**
   // Return the centroids
12 **return** Y

The algorithm consists of 4 main blocks:

- the **initialization block**, where all the centroids will receive a starting, random position (as per Lloyd's method);
- the **assignment block**, where the Euclidean distance between a point and all centroids is computed, for all centroids. The point will be assigned to a cluster depending on the following operation:

$$\underset{k}{\mathrm{argmin}} \left(\|\mu_k - x_i\| \right)$$

- the **update block**, where the position of the centroids is updated, and the new position of a centroid μ_k is equal to the mean of all the data points positions belonging to cluster k

B. Sequential Code Bottlenecks

For implementing the k-means algorithm, we will base all the codebase upon the project made from pro-

fessors Diego García-Álvarez and Arturo Gonzalez-Escribano from the University of Valladolid. The code shown in this subsection is taken from their project, although slightly adapted for giving enough context in the code snippets.

We described before the overall structure of the k-means algorithm: we will now proceed to examine its two main bottlenecks. As we can see from Algorithm 1, we have two main blocks that may cause performance issues: the **assignment block** and the **update block**.

The first **for** block in the **initialization step** does not represent a major bottleneck, since it just needs to assign a random location to each of the K centroids. It can be parallelized, but it won't help as much as parallelizing the two steps mentioned before.

The second **for** block represents the **assignment step**, which is, unlike the initialization step, computationally expensive: for each point, the algorithm will have to compute the euclidean distance (here onwards denotes as ℓ_2) between said point and all centroids $\mu_k \in C$, and select the lowest distance. This will determine the cluster of the point. In a C program, this may be accomplished with the following piece of code:

```
1  int cluster;
2  // For each point...
3  for(i = 0; i < points_number; i++) {
4      class = 1;
5      minDist = FLT_MAX;
6      // For each cluster...
7      for(j = 0; j < K; j++) {
8          // Compute the distance
9          dist = l2_norm(&data[i*samples], &
           centroids[j*samples], samples);
10
11         // If the distance is the lowest so far,
           replace it
12         if(dist < minDist) {
13             minDist = dist;
14             class = j+1;
15         }
16     }
17
18     // If the class is different from before, add
           a change to the counter
19     if(classMap[i] != class) {
20         changes++;
21     }
22
23     classMap[i]=class;
24 }
```

Notice the presence of the two nested **for** loops: sequentially, they would take a time of $O(|X| \cdot |C|)$, which may be optimized just by taking a simple single instruction multiple data approach (indeed,

with $m > 1$ different processes or threads, it would take a time of $O\left(\frac{|X| \cdot |C|}{m}\right)$ each, which is already better than the first option).

The third **for** loop represents the update step, which also is computationally expensive: we would need to perform the mean of the coordinates of all the points belonging to a cluster μ_k. This implies that all the coordinates of the points must be first summed, and then averaged on the number of points being classified to μ_k. An implementation in the C language would look like the following:

```
1  // For each point...
2  for (i = 0; i < lines; i++) {
3      point_class = classMap[i];
4      // Add 1 to the points classified for class k
5      pointsPerClass[point_class - 1] += 1;
6
7      // For each dimension...
8      for(j = 0; j < samples; j++) {
9          // ...add it to a table for summing and
           averaging
10         auxCentroids[(point_class - 1) * samples
           + j] += data[i * samples + j];
11     }
12 }
13
14 for (i = 0; i < K; i++) {
15     for (j = 0; j < samples; j++) {
16         // Average all dimensions
17         auxCentroids[i * samples + j] /=
           pointsPerClass[i];
18     }
19 }
20
21 maxDist = FLT_MIN;
22 for (i = 0; i < K; i++) {
23     // Compute the moving distance, as a
           convergence check
24     distCentroids[i] = euclideanDistance(&
           centroids[i * samples], &auxCentroids[i *
           samples], samples);
25     if (distCentroids[i] > maxDist) {
26         maxDist = distCentroids[i];
27     }
28 }
```

II. Parallelizing with MPI

III. Parallelizing with OpenMP

As a mean of enhancing the performance of the k-means algorithm, we made use of the multi-thread library OpenMP. OpenMP is a widely-used API for multi-platform shared-memory parallel programming. Our primary goal was to leverage multiple CPU cores to expedite the computationally intensive parts of the algorithm: the **cluster assignment** and **centroid update** steps. The number of threads can be optionally specified as a parameter; if not provided, a default value (i.e. 8 threads) is used for parallel execution. We now show how we approached

the parallelization of the program, and the advantages of taking certain design decisions.

A. OpenMP Implementation Approach

1. Cluster Assignment: the cluster assignment step is inherently parallelizable since each data point can be processed independently. We made use of the `#pragma omp parallel for` directive to distribute the loop iterations across multiple threads. Said directive has been accompanied by the following clauses:

- `reduction(+ : changes)`: the **reduction clause** safely aggregates the total number of cluster reassignments across threads. This avoids race conditions by creating thread-private copies of `changes` that are combined after the loop;
- `schedule(dynamic, 16)`: the **dynamic scheduling** clause improves load balancing by allowing threads to dynamically claim "chunks" of 16 iterations.

2. Centroid Update: the centroid update step involves two parallel phases: gathering cluster sums and averaging the centroids coordinates. Here, we addressed potential race conditions through privatization and synchronization:

- **Privatization with Local Buffers**: Each thread maintains private copies of two buffers, namely `local_pointsPerClass` (count of points assigned to the cluster) and `local_auxCentroids` (cumulative sum of data points for each cluster). This eliminates races during local accumulation, as threads operate on isolated data;
- **Critical Section for Global Aggregation**: After local accumulation, a `#pragma omp critical` region safely merges thread-local results into global `pointsPerClass` and `auxCentroids`. While critical sections incur synchronization overhead, they are used sparingly here once per thread in order to minimize contention;
- **Parallel Mean Calculation**: The final centroid normalization loop (`#pragma omp parallel for`) parallelizes the division by `pointsPerClass`, avoiding race conditions since each cluster is processed independently.

An alternative approach that we experimented with uses the reduction feature to eliminate the need for explicit synchronization.

```
1  #pragma omp parallel for reduction(+ :
       pointsPerClass[ : K], auxCentroids[ : K *
       samples])
2          for (i = 0; i < lines; i++)
3          {
4              int class = classMap[i] – 1;
5              pointsPerClass[class]++;
6              for (int j = 0; j < samples; j++)
7              {
8                  auxCentroids[class * samples + j]
       += data[i * samples + j];
9              }
10         }
```

The reduction clause automates the process of merging thread-local results into shared variables. By reducing both `pointsPerClass` and `auxCentroids`, the need for a critical section is removed, potentially lowering synchronization overhead. However, we chose the original approach with explicit privatization and critical regions to demonstrate proficiency with OpenMP synchronization and race condition management.

The maximum centroid displacement (`maxDist`) is computed serially. While this loop could be parallelized using the `#pragma omp parallel for reduction(max : maxDist)` directive, we retained the serial approach for simplicity, as its computational cost is negligible compared to other steps and because the number of clusters k is typically small, the overhead associated with parallelization outweighs the potential performance gains.

B. Solutions Implemented

- **Reduction for Scalars**: The `changes` variable uses a `reduction` clause, which is more efficient than atomic operations for scalar summation. OpenMP handles private copies and post-loop merging automatically.
- **Privatization and Critical Sections**: For array/matrix updates (i.e., `auxCentroids`), thread-local buffers reduce the need for fine-grained synchronization. The critical section ensures safe aggregation with minimal overhead, as it is invoked only once per thread.
- **Atomic Operations vs. Critical Regions**: While atomic operations (i.e., `#pragma omp atomic`) could replace the critical section for scalar increments (`pointsPerClass`), they would be inefficient for matrix additions (`auxCentroids`) due to repeated fine-grained locking. Privatization strikes a better balance between correctness and performance.

C. Performance Considerations

- **Critical Section Overhead**: The single critical section per thread during centroid aggregation has negligible cost compared to the computational work, as merging local buffers is a minor operation.
- **Memory Efficiency**: Privatizing per thread the `local_auxCentroids` buffer increases memory usage proportionally to the number of threads. However, this is manageable for moderate thread counts and cluster sizes.
- **Dynamic Scheduling Trade-off**: While dynamic scheduling improves load balancing, it introduces overhead for chunk management. A chunk size of 16 was empirically chosen to balance parallelism and scheduling latency.

IV. Parallelizing with CUDA

V. Interlacing Multi-processing with Multi-threading

So far, we implemented various solutions for our programs, which employed either multi-processing or multi-threading parallelism techniques, without using both approaches at the same time. However, these techniques are not mutually exclusive, and can be mixed together in order to achieve better performances. In fact, in high performance computing tasks, multi-process techniques are used within clusters to connect nodes and coordinate them, while multi-threading are used for performing computations, given the directives of the master process(es). In this section we will show how the two multi-threading approaches that we previously considered (namely, OpenMP and CUDA) can be combined with the famous multi-processing library MPI.

A. MPI and OpenMP

The hybrid MPI + OpenMP implementation of the *k*-means clustering algorithm leverages both distributed and shared memory parallelism to enhance performance and scalability. Our parallelization strategy can be described with three main key points:

Two-Level Parallelism

- MPI distributes data across processes (coarse-grained parallelism) with `MPI_Scatterv`, a vector scattering collective, while OpenMP parallelizes loops within each process (allowing for a much fine-grained parallelism);

- Thread safety: all the processes are initialized with `MPI_THREAD_FUNNELED`, so that to restrict MPI calls to the main thread.

Asynchronous Communication

The non-blocking reduce operation `MPI_Ireduce` aggregates cluster reassignment counts while centroid updates proceed, allowing to overlap communication and computation.

Hybrid Reductions

- OpenMP sets thread-local buffers in order to avoid intra-node races. More specifically, the `local_pointsPerClass` and the `local_auxCentroids` buffers;
- Global aggregations via `MPI_Allreduce` ensures consistent centroid states across nodes.

```
1 #pragma omp parallel  // Thread-local
      accumulation
2 {
3     int *local_pointsPerClass = calloc(K, sizeof(
      int));
4     float *local_auxCentroids = (float *)calloc(K
       * samples, sizeof(float));
5     // ...
6 }
7
8 // Global sync
9 MPI_Allreduce(MPI_IN_PLACE, pointsPerClass, K,
      MPI_INT, MPI_SUM, MPI_COMM_WORLD);
10 MPI_Allreduce(MPI_IN_PLACE, auxCentroids, K *
      samples, MPI_FLOAT, MPI_SUM, MPI_COMM_WORLD);
```

A major concern while designing this application was also ensuring that an efficient **data distribution** system was in place. We show, in the following two key points, how we tackled this problem:

Balanced Workloads:

- The `MPI_Scatterv` collective allows to split data into chunks, adjusted via two parameters: `sendcounts` and `displs`. Said collective is capable of handling uneven divisions, ensuring balanced workloads even when the total data points (`lines`) are not evenly divisible by the number of processes.
- Each process computes locally the updates for a subset of centroids. These partial updates are then gathered from all processes and redistributed to ensure that every process has the complete and updated set of centroids for the next iteration, using `MPI_Allgatherv`. We report here the full collective call:

```
1 MPI_Allgatherv(local_centroids, local_k *
      samples, MPI_FLOAT, centroids,
      centroid_sendcounts, centroid_displs,
      MPI_FLOAT, MPI_COMM_WORLD);
```

Final Result Aggregation

After convergence, the `MPI_Gatherv` collective reconstructs the global cluster assignments on the root process. Each process sends its own `local_classMap` (containing classifications for its local data subset), and the root process assembles them into `classMap` using the precomputed `recvcounts` and `rdispls` offsets. Below the code snippet, which shows how the `MPI_Gatherv` call was assembled, we give more precise definitions regarding the two offsets aforementioned:

```
1 // Gather local classifications into root's
      classMap
2 MPI_Gatherv(local_classMap, local_lines, MPI_INT,
       classMap, recvcounts, rdispls, MPI_INT, 0,
      MPI_COMM_WORLD);
```

- `recvcounts`: array specifying the number of elements received from each process (matches `sendcounts` from initial scatter);
- `rdispls`: offsets in `classMap` where each process' data is stored.

These offsets ensures efficient reconstruction of the global result without requiring intermediate storage on worker processes.

Memory Efficiency:
- Each process stores only its local data (`local_points`) and centroid subset (`local_centroids`).
- Thread-local buffers minimize shared memory contention without atomic operations.

How did we approach **synchronization and termination**?

Convergence Check:
- OpenMP's `reduction(max:local_maxDist)` computes per-process centroid displacements.
- `MPI_Reduce` finds the global maximum displacement for termination decisions.

Termination Broadcast

Root process evaluates conditions and broadcasts `terminate` flag via `MPI_Bcast`, ensuring synchronized loop exits.

B. MPI and CUDA

VI. PERFORMANCE ANALYSIS

VII. CONCLUSIONS